

## **Phase-4: Development Part II**

**Title:** AI Based Diabetes Prediction System

**Dataset link:** <https://www.kaggle.com/datasets/mathchi/diabetes-data-set>

**Software used:** google colab

### **Introduction:**

Diabetes is a health condition that affects how your body turns food into energy. Most of the food you eat is broken down into sugar (also called glucose) and released into your bloodstream. When your blood sugar goes up, it signals your pancreas to release insulin.

### **Previous step:**

- Installing Libraries
- Importing data
- Displaying data
- Data Preprocessing
- Missing value analysis

**Next:**

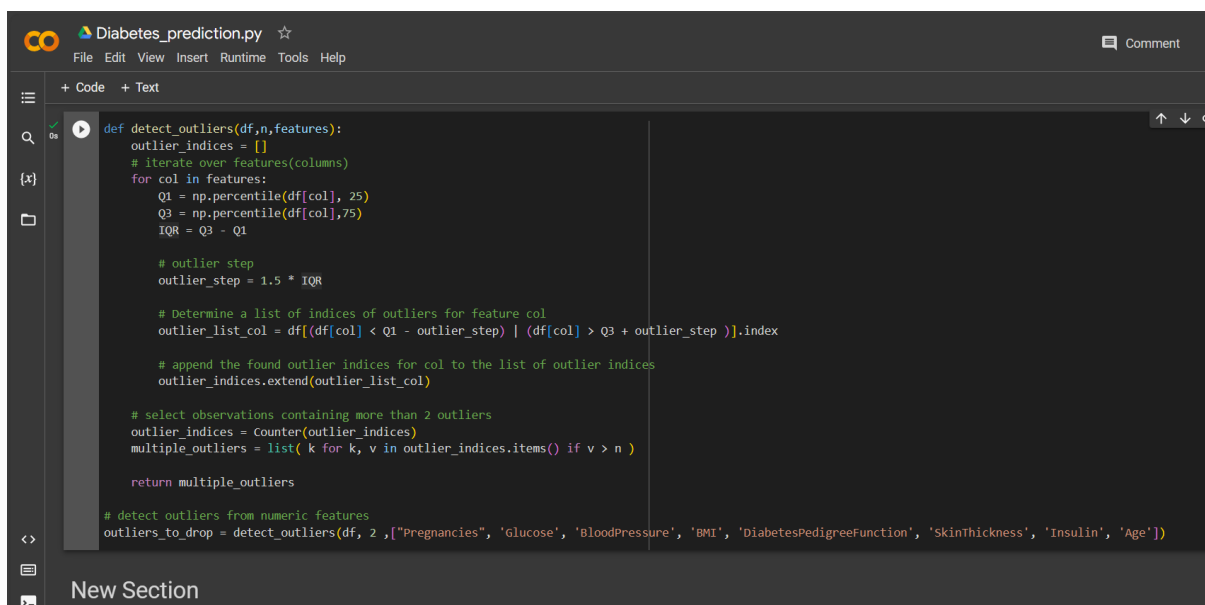
## Feature Engineering:

Till now, i explored the dataset, did missing value corrections and data visualization. Next, i have started feature engineering. Feature engineering is useful to improve the performance of machine learning algorithms and is often considered as applied machine learning.

Selecting the important features and reducing the size of the feature set makes computation in machine learning and data analytic algorithms more feasible.

## Outlier Detection:

In this part i removed all the records outlined in dataset. Outliers impacts Model accuracy. I used *Tukey Method* used for outlier detection.



```
Diabetes_prediction.py
File Edit View Insert Runtime Tools Help

+ Code + Text

def detect_outliers(df,n,features):
    outlier_indices = []
    # iterate over features(columns)
    for col in features:
        Q1 = np.percentile(df[col], 25)
        Q3 = np.percentile(df[col],75)
        IQR = Q3 - Q1

        # outlier step
        outlier_step = 1.5 * IQR

        # Determine a list of indices of outliers for feature col
        outlier_list_col = df[(df[col] < Q1 - outlier_step) | (df[col] > Q3 + outlier_step )].index

        # append the found outlier indices for col to the list of outlier indices
        outlier_indices.extend(outlier_list_col)

    # select observations containing more than 2 outliers
    outlier_indices = Counter(outlier_indices)
    multiple_outliers = list( k for k, v in outlier_indices.items() if v > n )

    return multiple_outliers

# detect outliers from numeric features
outliers_to_drop = detect_outliers(df, 2 ,["Pregnancies", 'Glucose', 'BloodPressure', 'BMI', 'DiabetesPedigreeFunction', 'SkinThickness', 'Insulin', 'Age'])
```

Here, I find outliers from all the features such as Pregnancies, Glucose, BloodPressure, BMI, DiabetesPedigreeFunction, SkinThickness, Insulin, and Age.

```
df.drop(df.loc[outliers_to_drop].index, inplace=True)
print(df)
```

I have successfully removed all outliers from dataset now. The next step is to split the dataset in train and test and proceed the modeling.

## Modeling:

In this sections, i tried different models and compare the accuracy for each. Then, i performed Hyperparameter Tuning on Models that has high accuracy.

Before i split the dataset i need to transform the data into quantile using `sklearn.preprocessing`.

```
# Data Transformation
q = QuantileTransformer()
X = q.fit_transform(df)
transformedDF = q.transform(X)
transformedDF = pd.DataFrame(X)
transformedDF.columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome']
# Show top 5 rows
transformedDF.head()
```

/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/\_data.py:2627: UserWarning: n\_quantiles (1000) is greater than the total number of samples  
warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but QuantileTransformer was fitted with fea  
warnings.warn(

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	0.747718	0.810300	0.516949	0.801825	0.000000	0.591265	0.750978	0.889831	1.0
1	0.232725	0.097784	0.336375	0.644720	0.000000	0.227510	0.475880	0.558670	0.0
2	0.863755	0.956975	0.279009	0.000000	0.000000	0.091917	0.782269	0.585398	1.0
3	0.232725	0.131030	0.336375	0.505867	0.662973	0.298566	0.106258	0.000000	0.0
4	0.000000	0.721643	0.050847	0.801825	0.834420	0.926988	0.997392	0.606258	1.0

## Data Splitting:

Next, i split data in test and train dataset. Train dataset will be used in Model training and evaluation and test dataset will be used in prediction. Before i predict the test data, i performed cross validation for various models.

```
features = df.drop(["Outcome"], axis=1)
labels = df["Outcome"]
x_train, x_test, y_train, y_test = train_test_split(features, labels, test_size=0.30, random_state=7)
```

Above code splits dataset into train (70%) and test (30%) dataset.

## Cross Validate Models:

```
def evaluate_model(models):
    """
    Takes a list of models and returns chart of cross validation scores using mean accuracy
    """
    # Cross validate model with Kfold stratified cross val
    kfold = StratifiedKFold(n_splits = 10)

    result = []
    for model in models :
        result.append(cross_val_score(estimator = model, X = x_train, y = y_train, scoring = "accuracy", cv = kfold, n_jobs=4))
```

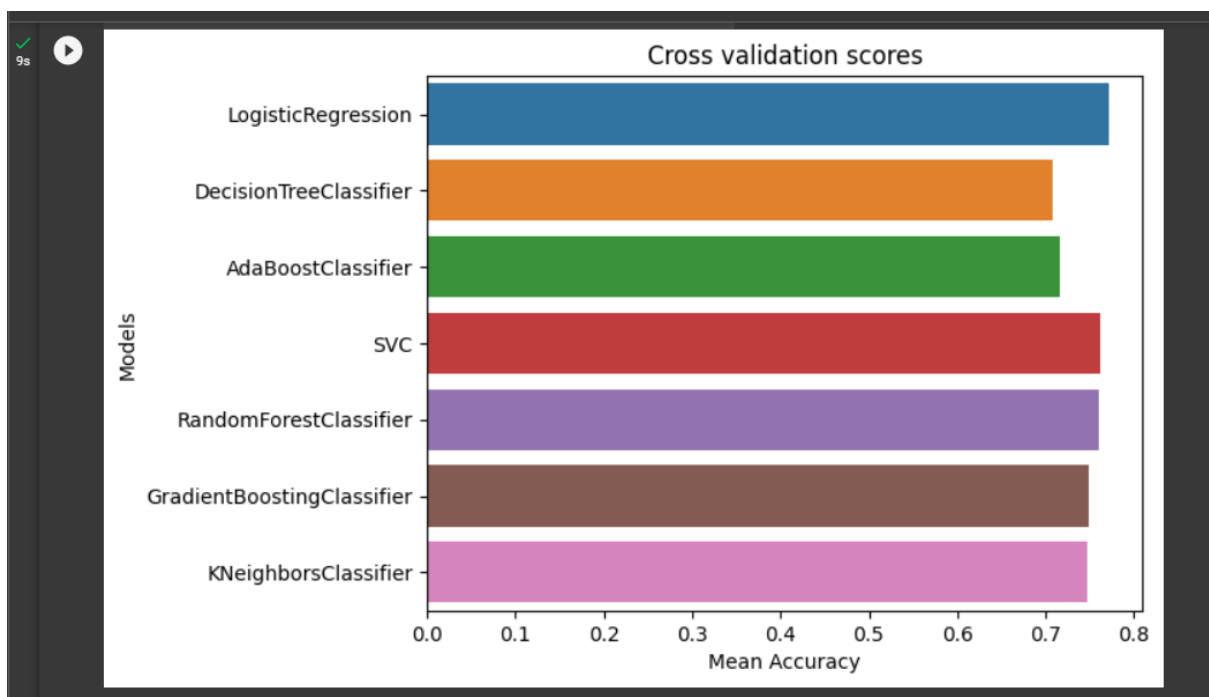
```
result_df = pd.DataFrame({
    "CrossValMeans":cv_means,
    "CrossValerrors": cv_std,
    "Models":[
        "LogisticRegression",
        "DecisionTreeClassifier",
        "AdaBoostClassifier",
        "SVC",
        "RandomForestClassifier",
        "GradientBoostingClassifier",
        "KNeighborsClassifier"
    ]
})

# Generate chart
bar = sns.barplot(x = "CrossValMeans", y = "Models", data = result_df, orient = "h")
bar.set_xlabel("Mean Accuracy")
bar.set_title("Cross validation scores")
return result_df
```

*Method `evaluate\_model` takes a list of models and returns chart of cross validation scores using mean accuracy.*

```
9s # Modeling step Test differents algorithms
random_state = 30
models = [
    LogisticRegression(random_state = random_state, solver='liblinear'),
    DecisionTreeClassifier(random_state = random_state),
    AdaBoostClassifier(DecisionTreeClassifier(random_state = random_state), random_state = random_state, learning_rate = 0.2),
    SVC(random_state = random_state),
    RandomForestClassifier(random_state = random_state),
    GradientBoostingClassifier(random_state = random_state),
    KNeighborsClassifier(),
]
evaluate_model(models)
```

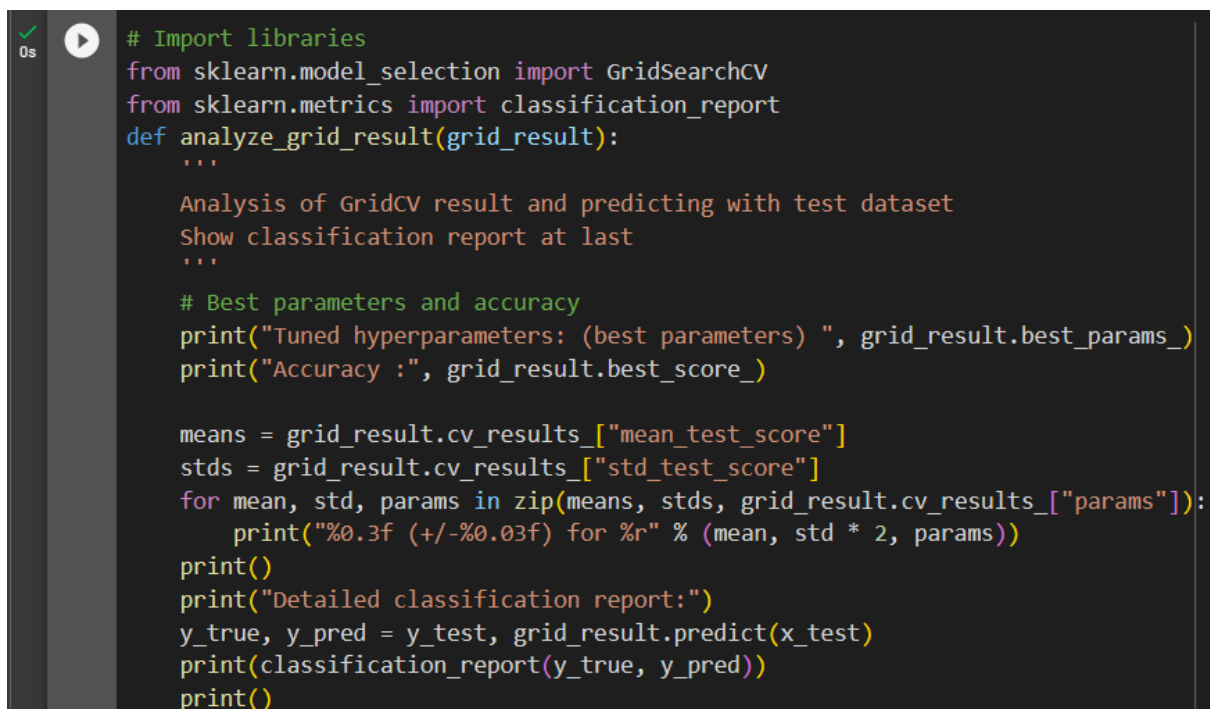
	CrossValMeans	CrossValerrors	Models
0	0.770964	0.058524	LogisticRegression
1	0.707687	0.065109	DecisionTreeClassifier
2	0.717016	0.061262	AdaBoostClassifier
3	0.761670	0.033724	SVC
4	0.759748	0.055512	RandomForestClassifier
5	0.748532	0.065303	GradientBoostingClassifier
6	0.746506	0.054171	KNeighborsClassifier



As per above observation, I found that Logistic Regression model has more accuracy. next, I will do hyperparameter tuning on model.

## Hyperparameter Tuning:

Hyperparameter tuning is choosing a set of optimal hyperparameters for a learning algorithm. A hyperparameter is a model argument whose value is set before the learning process begins. The key to machine learning algorithms is hyperparameter tuning.

A screenshot of a code editor with a dark background. On the left, there is a vertical sidebar with a green checkmark icon and the text '0s'. The main area contains Python code for hyperparameter tuning. The code imports GridSearchCV and classification\_report from sklearn, defines a function analyze\_grid\_result, and then uses it to tune a model. The function prints the best parameters, accuracy, and a detailed classification report.

```
# Import libraries
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report
def analyze_grid_result(grid_result):
    """
    Analysis of GridCV result and predicting with test dataset
    Show classification report at last
    """

    # Best parameters and accuracy
    print("Tuned hyperparameters: (best parameters) ", grid_result.best_params_)
    print("Accuracy :", grid_result.best_score_)

    means = grid_result.cv_results_["mean_test_score"]
    stds = grid_result.cv_results_["std_test_score"]
    for mean, std, params in zip(means, stds, grid_result.cv_results_["params"]):
        print("%.3f (+/-%.03f) for %r" % (mean, std * 2, params))
    print()
    print("Detailed classification report:")
    y_true, y_pred = y_test, grid_result.predict(x_test)
    print(classification_report(y_true, y_pred))
    print()
```

First of all i have imported GridSearchCV and classification\_report from sklearn library. Then, i have defined `analyze\_grid\_result` method which will show prediction result. I called this method for each Model used in SearchCV. In next step, i will perform tuning for each model.

## Logistic Regression:

```
11s # Define models and parameters for LogisticRegression
model = LogisticRegression(solver='liblinear')
solvers = ['newton-cg', 'liblinear']
penalty = ['l2']
c_values = [100, 10, 1.0, 0.1, 0.01]
# Define grid search
grid = dict(solver = solvers, penalty = penalty, C = c_values)
cv = StratifiedKFold(n_splits = 50, random_state = 1, shuffle = True)
grid_search = GridSearchCV(estimator = model, param_grid = grid, cv = cv, scoring = 'accuracy', error_score = 0)
logi_result = grid_search.fit(x_train, y_train)
# Logistic Regression Hyperparameter Result
analyze_grid_result(logi_result)
```

## Output:

```
🔍 Tuned hyperparameters: (best parameters) {'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}
Accuracy : 0.774909090909091
0.773 (+/-0.241) for {'C': 100, 'penalty': 'l2', 'solver': 'newton-cg'}
0.773 (+/-0.241) for {'C': 100, 'penalty': 'l2', 'solver': 'liblinear'}
0.773 (+/-0.241) for {'C': 10, 'penalty': 'l2', 'solver': 'newton-cg'}
0.775 (+/-0.226) for {'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}
0.773 (+/-0.240) for {'C': 1.0, 'penalty': 'l2', 'solver': 'newton-cg'}
0.773 (+/-0.224) for {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}
0.773 (+/-0.242) for {'C': 0.1, 'penalty': 'l2', 'solver': 'newton-cg'}
0.720 (+/-0.225) for {'C': 0.1, 'penalty': 'l2', 'solver': 'liblinear'}
0.764 (+/-0.245) for {'C': 0.01, 'penalty': 'l2', 'solver': 'newton-cg'}
0.687 (+/-0.256) for {'C': 0.01, 'penalty': 'l2', 'solver': 'liblinear'}

Detailed classification report:

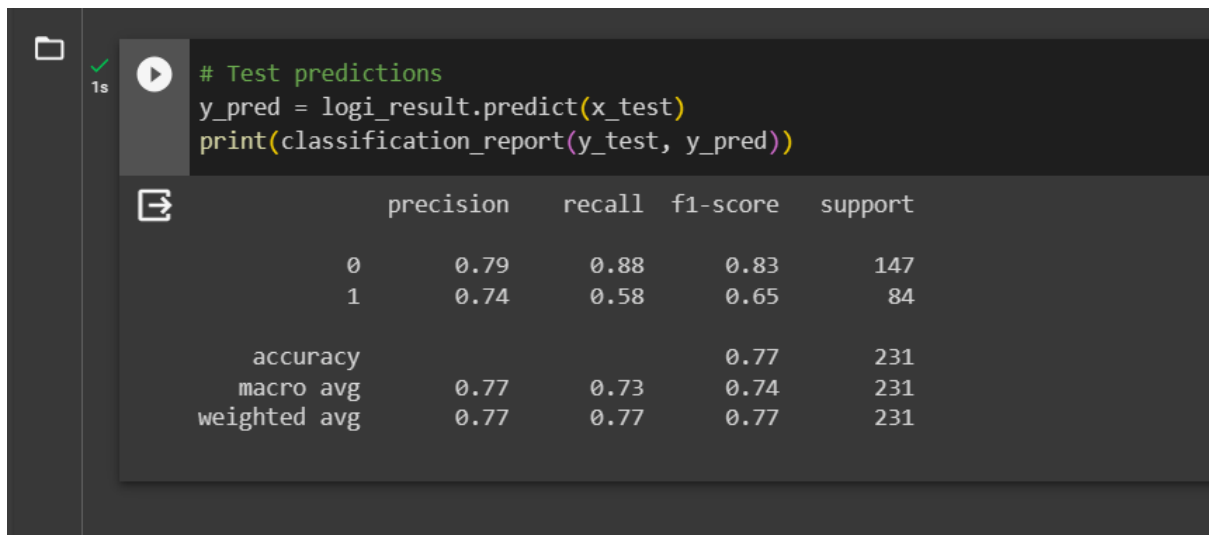
```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	147
1	0.74	0.58	0.65	84
accuracy			0.77	231
macro avg	0.77	0.73	0.74	231
weighted avg	0.77	0.77	0.77	231

As per my observation, in LogisticRegression it returned best score 0.78 with `{'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}` parameters.

## Prediction:

Till now, i worked on Feature Engineering, Cross Validation of Models, and Hyperparameter Tuning and find the best working Model for my dataset. Next, I did prediction from my test dataset and storing the result in CSV.



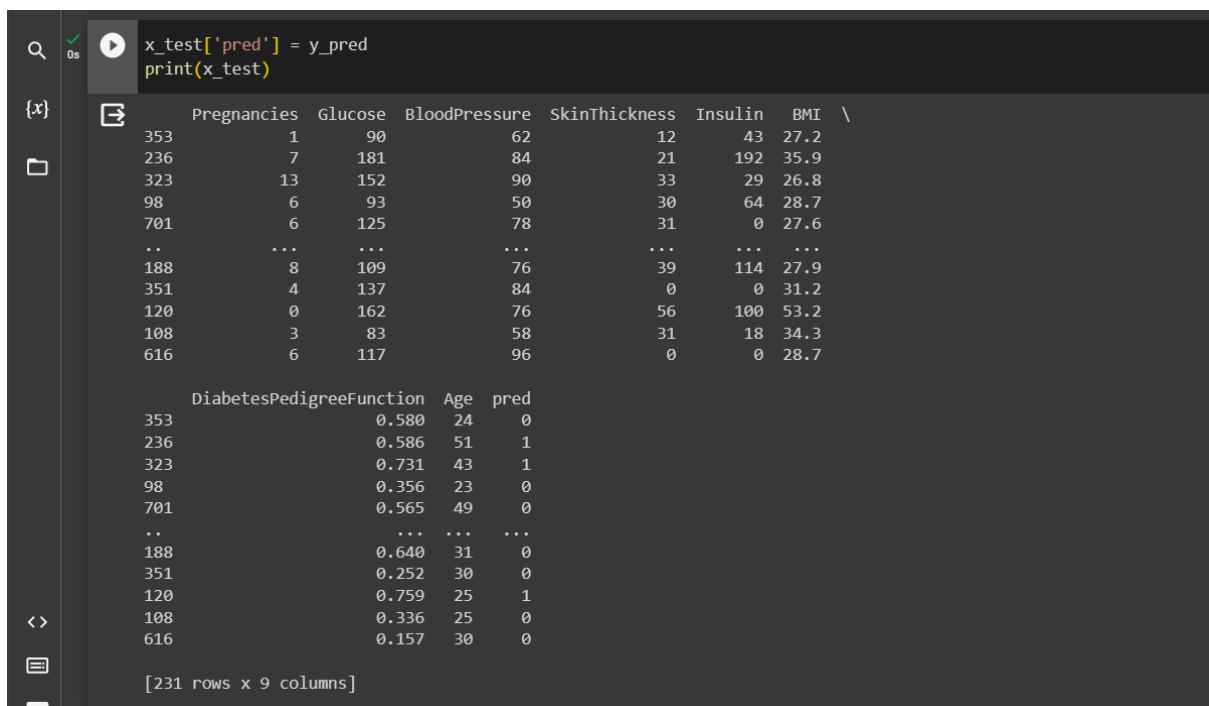
The screenshot shows a Jupyter Notebook cell with the following code and output:

```
# Test predictions
y_pred = logi_result.predict(x_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	147
1	0.74	0.58	0.65	84
accuracy			0.77	231
macro avg	0.77	0.73	0.74	231
weighted avg	0.77	0.77	0.77	231

Finally append new feature column in test dataset called Prediction and print the dataset.

## Final output:



The screenshot shows a Jupyter Notebook cell with the following code and output:

```
x_test['pred'] = y_pred
print(x_test)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
353	1	90	62	12	43	27.2	
236	7	181	84	21	192	35.9	
323	13	152	90	33	29	26.8	
98	6	93	50	30	64	28.7	
701	6	125	78	31	0	27.6	
...	...	...	...	...	...	...	
188	8	109	76	39	114	27.9	
351	4	137	84	0	0	31.2	
120	0	162	76	56	100	53.2	
108	3	83	58	31	18	34.3	
616	6	117	96	0	0	28.7	
	DiabetesPedigreeFunction	Age	pred				
353	0.580	24	0				
236	0.586	51	1				
323	0.731	43	1				
98	0.356	23	0				
701	0.565	49	0				
...	...	...	...				
188	0.640	31	0				
351	0.252	30	0				
120	0.759	25	1				
108	0.336	25	0				
616	0.157	30	0				

[231 rows x 9 columns]



**Conclusion:**

1. Diabetes is one of the risks during Pregnancy. It has to be treated to avoid complications.
2. BMI index can help to avoid complications of diabetes a way before
3. Diabetes starts showing in age of 35 – 40 and increases with person age.